NUBRA · PRODUCT MANAGEMENT INTERN CASE · PROBLEM 03, POST-TRADE REVIEW

TradeLens: the monthly review that names the habits costing you money

Bhavya Mohnot · October 2026

My take on the problem

Traders don't skip reviews out of laziness; reviewing is hard. A tradebook is hundreds of buy/sell rows, and a P&L report says *how much* you lost, not *what to change*. The losses that matter are repeated behaviours: the instrument that never works for you, the expiry-day option punt, the trade taken ten minutes after a loss. Each is invisible one trade at a time and obvious across a month.

87.7%
of individual F&O traders lost money in FY26

₹91,685 cr
lost by individual traders in total

~90%
of those who lost two years running, and kept trading, lost again

People keep trading; they don't keep learning.

> So the job is not "more analytics". It is to take a month of raw order history, name the **two or three habits that cost the most rupees**, prove each with the actual trades, and hand back **one rule to try next month**.

Why Nubra should solve it

- **It already has the data.** Nothing to export or upload; the review can arrive on the 1st of every month.
- **The right users.** Nubra is built for active traders, scalpers and options strategists, which is where habits compound fastest.
- **It can close the loop.** A journal app can only tell you about a habit. Nubra can turn it into a guardrail at the moment of decision: an expiry-day nudge on the order ticket, a cool-off after three losses, or pausing an instrument on a smart watchlist.
- **Retention.** Traders who blow up their accounts churn. Helping them lose less is good for them and for lifetime value.

The solution space

OPTION	WHY NOT (OR WHY)
P&L dashboard (Console-style)	Answers "how much", not "what to change"; most traders never open it
Trade journal (TradeZella, Tradervue)	Needs manual tagging; too much effort for most retail traders
AI chat over your trades	Flexible but hard to verify; a confident wrong rupee figure does real damage
Opinionated monthly review (chosen)	Deterministic checks, ranked by ₹, at most three, each with evidence and a rule

1. **Rank by marginal ₹.** If BANKNIFTY losses and closing-hour losses are the same trades, they count once, and the closing hour is cleared.
2. **Don't invent problems.** A habit must cover at least 5% of the month's losses and be statistically worse than the trader's other trades. Everything checked and cleared is shown, for example "your quick re-entries made +₹650".
3. **Plain words, no motive-guessing.** It says "jumping back in within 15 minutes of a loss", never "revenge trading".

In scope (v1)

- One month of NSE equity/F&O order history: a broker tradebook CSV/XLSX, or a closed-trade sheet.
- Round-trip rebuild from fills (scale-ins, partial exits, shorts), with charges estimated when the file has none.
- Nine habit checks: an instrument that keeps losing · trading on after three losses · quick re-entry after a loss · overtrading · a losing time of day · expiry-day options · holding losers longer than winners · sizing up after a loss · charges eating small trades.
- A one-rule commitment, checked against next month's review.
- Runs entirely in the browser; the file never leaves the device.

Out of scope

- Live positions and margin (problem 02).
- Broker API sync.
- Grouping multi-leg strategies (each leg is reviewed on its own).
- Intraday price paths, so no "you exited too early" claims.
- Tax P&L, guessing emotions, free-form AI chat.

How I'd know it works

- **Adoption:** the share of monthly active traders who open the review, and the share who commit to a rule.
- **Outcome (north star):** the ₹ cost of the committed habit the following month against its baseline. Target: at least a 20% drop for half of committers.
- **Trust:** the "this is wrong" rate per habit card; a high rate means re-tuning or removing that check.

Figures: SEBI study of individual traders in equity derivatives, FY26 (Aug 2026). Live app: trade-lens-omega.vercel.app · Prompts p.2–3 · Evals p.4 · App p.5

ARTIFACT 02 · PROMPTS

The prompt trail, from first try to shipped app

Round 1 was built quickly with ChatGPT. Round 2 audited it against the brief and rebuilt it with Claude. The prompts are quoted as sent, typos included, each with what came back.

didn't work · failed, then fixed · worked

Round 1 · ChatGPT

✓ PROMPT 1

```
which of these can we make immediately
```

Compared the four problems and recommended **#03 Post-trade review**: clearest problem, fewest dependencies, testable.

✕ PROMPT 2

```
lets quickly build an agentic ai workflow with a web app where chatbot type interface, chat history and so on
is there where the traders can upload their trade book and the agentic workflow will analyse it for them ,
behavioural patterns, wins losses and so on so forth
```

A chat-style dashboard: upload, saved history, win/loss/fee metrics. The AI was never connected; the "chat" matched keywords. **It described a technology, not the trader's job.**

✕ PROMPT 3

```
deploy on vercel real quick and no login required. and lets test it out real quick
```

Deployed, but only the compiled output was uploaded. **The source code was never kept**, which mattered later.

✓ PROMPTS 4–5

```
make it more open ,, the excel upload
```

```
even more, i meant the pnl and other columns requirement sounds big, or have an example created for me real
quick so i can test
```

Excel import, plus **TradeLens-example.xlsx**: 60 synthetic trades with three planted patterns. It became the first test file for Round 2.

Round 2 · Claude

✓ PROMPT 6 · AUDIT

```
I have this project: [Round 1 app URL]
Can you access this project and then do an analysis. This project should be suitable for this: "[project
folder]". It should match all the constraints and requirements in teh Nubra Product Intern Assignment PDF.
Also, it should behave appropriately to the Example data in the directory. Also check that the project has all
the requirements mentioned in the JD.
```

The audit found that Round 1 **ranked the wrong habit first**: fees ₹2,100 above BANKNIFTY ₹9,100. It **flagged re-entries that made money** (+₹650), answered "worst hour" with the **best** one, **couldn't read order history**, and had no one-pager or evals.

✕ PROMPTS 7–8 · RECOVER THE CODE

```
can you access the source code of this website
```

```
i am on the source tab, check in this screenshot: [Vercel → Source tab]
```

Dead end: no source maps, and Vercel held only the built files. Nothing to fix in place.

Continued overleaf. Prompts about account setup and file formatting are left out.

The prompt trail, continued

Round 2 rebuilt the app from scratch, then tested, fixed and shipped it.

✓ PROMPT 9 · REBUILD

```
Okay, so scratch that, just do this: Create for me everything this project needs from scratch, all the
artifacts (submissions), meeting all the requirements with fresh source code. Keep the source code local for
now, we will be directly uploading it to vercel.
```

Claude reframed the product around one outcome: **at most three habits, ranked by marginal ₹, each proven with trades and turned into one rule**. It then wrote the engine: it reads broker tradebooks, rebuilds round trips, estimates NSE charges, and runs nine habit checks. On the brief's example file it named **BANKNIFTY (₹9,100)** and cleared the profitable re-entries.

✓ PROMPT 10 · DATA

```
can you also generate a better/improved example dataset for this project, which meets all the JD and PDF
requirements.
```

A realistic **September 2026 Zerodha-style tradebook**: 388 buy/sell fills across NIFTY weeklies, BANKNIFTY, RELIANCE futures and equity. It has three planted habits, decoys that must not be flagged, and an answer key computed independently of the app.

✓ PROMPT 11 · INTERFACE (EXCERPT)

```
…single narrative column (max-width ~760px). It must NOT feel like a dashboard. Order matters: …
…never use the word "revenge", never say "you were emotional"…
…must work at 375px wide…
```

Built and type-checked; works at desktop and phone width.

↻ PROMPT 12 · EVALS (EXCERPT)

```
…Planting functions, one per habit key above, each with a `strength` parameter…
…recall@3 ≥ 90% and top-1 ≥ 80% for every habit type; clean-control FP rate ≤ 10%…
…Do not loosen thresholds or assertions to make things pass.
```

First run failed the ship bar. Expiry-day losses were hidden under "NIFTY keeps losing" for 25% of traders, and one test expected the wrong number (a 20-trade losing streak; it is 21). False alarms on clean traders sat right at the 10% limit.

Fixed: prefer the narrower habit when it explains most of the cost (75% → 100%), and use a stricter bar for checks that test many variants (false alarms 10% → 0%). The test was corrected, with the reason recorded.

↻ FURTHER TESTING

The evals were rerun on seeds never used for tuning, and each habit's ₹ cost was compared with its own evidence. **Both failed:** the "charges eating small trades" check dropped to 75–80% on unseen data, and "holding losers" claimed far more ₹ than its averages supported.

Fixed: the planted charges sat exactly on the 5% materiality line, so they were rescaled. Holding losers is now measured only after higher habits claim their trades. Final result on unseen data: **95% recall, 3% false alarms**.

↻ PROMPT 13

```
show me the app
```

The main sample button failed: "We could not load the sample file", and one data note appeared twice.

Fixed: the dev server hadn't picked up the new sample (the production build had it), and the duplicate note was removed.

✓ PROMPT 14 · SHIP (EXCERPT)

```
push the updates to github yourself … Once you've done that, confirm if everything is reflected on the hosted
vercel url
```

Shipped. It is live at trade-lens-omega.vercel.app. Both samples were checked on the live site at desktop and phone width. The September sample shows **three habits costing ₹57,978**, and on the brief's file **BANKNIFTY ₹9,100**. 91 unit tests and 22 fixed eval cases pass.

ARTIFACT 03 · EVALS

How I decided it was good enough to ship

Every ₹ figure comes from deterministic code, so every claim can be tested. A single command, `npm run eval`, runs everything and fails if the ship bar isn't met.

Four layers of checks

LAYER	WHAT IT PROVES	RESULT
Unit tests (91)	NSE symbol parsing incl. weekly/monthly expiries; timestamp formats; round-trip rebuild for scale-ins, partial exits, shorts, flips and open positions	91 / 91 pass
Fixed cases (22)	The brief's example file to the rupee; the September sample against its answer key; one case per habit; a clean trader; a symbol-vs-time confound; 8 messy-format cases (net-only P&L, no timestamps, DD-MM dates with ₹ and brackets, BOM/CRLF, bad rows, no price column, empty files, open position)	22 / 22 pass
Seeded sweeps	Each habit planted into 20 synthetic traders: is it among the 3 shown (recall@3), and is it #1? Plus 30 clean traders with nothing planted (false alarms)	See below
Copy & consistency	On all 350 reviews: titles ≤60 chars, no judgemental words ("revenge", "gamble"…), every habit has ₹, trades and a rule, ranks sorted by cost, all ₹ reconcile	0 violations

Ship bar, on tuning seeds and on unseen seeds

METRIC	BAR	TUNING SEEDS	HELD-OUT SEEDS (NEVER TUNED ON)
Planted habit is shown (recall@3), worst habit type	≥ 90%	100%	95%
Planted habit ranked #1, worst habit type	≥ 80%	95%	85%
False alarms on clean traders	≤ 10%	0%	3%
Fixed cases · copy checks	all	pass	pass

Thresholds were tuned on seeds 1–30 and 1001–1030, then checked once on seeds 2001–2030. A habit at half strength is mostly *not* reported (0–20% at half strength); that is by design: silence beats a false alarm.

Acceptance on the example data

Brief's example file (60 trades)

- Net −₹5,090 (gross −₹2,990, charges ₹2,100) ✓
- One habit: **BANKNIFTY kept losing**, ₹9,100 ✓
- Re-entries cleared: they made +₹650 ✓
- Closing hour cleared once BANKNIFTY is set aside ✓
- Weekend-dated trades flagged as a data warning ✓

September 2026 tradebook (388 fills)

- 158 round trips, 1 open position set aside ✓
- #1 Expiry-day options ₹30,131 · #2 Quick re-entry ₹15,932 · #3 Holding losers ₹11,916, matching the answer key ✓
- No decoy flagged: profitable opening half hour, RELIANCE ✓

What the evals caught, and what changed

1. **Expiry-day losses were hidden under "NIFTY keeps losing"** in 25% of traders. The engine now prefers the narrower habit when it explains most of the cost (75% → 100%).
2. **10% false alarms on clean traders**, all from checks that test many variants at once (5 time windows, every instrument). A stricter bar for those cut it to 0% (3% unseen).
3. **"Holding losers" overstated its cost**, then was masked by other habits' fast losses. It is now measured only after higher habits claim their trades, so its evidence and ₹ agree.
4. **Twice the test was wrong, not the app:** the example's longest losing streak is 21, not 20; and a planted habit sat exactly on the 5% materiality line. Both were fixed in the tests, with reasons recorded.

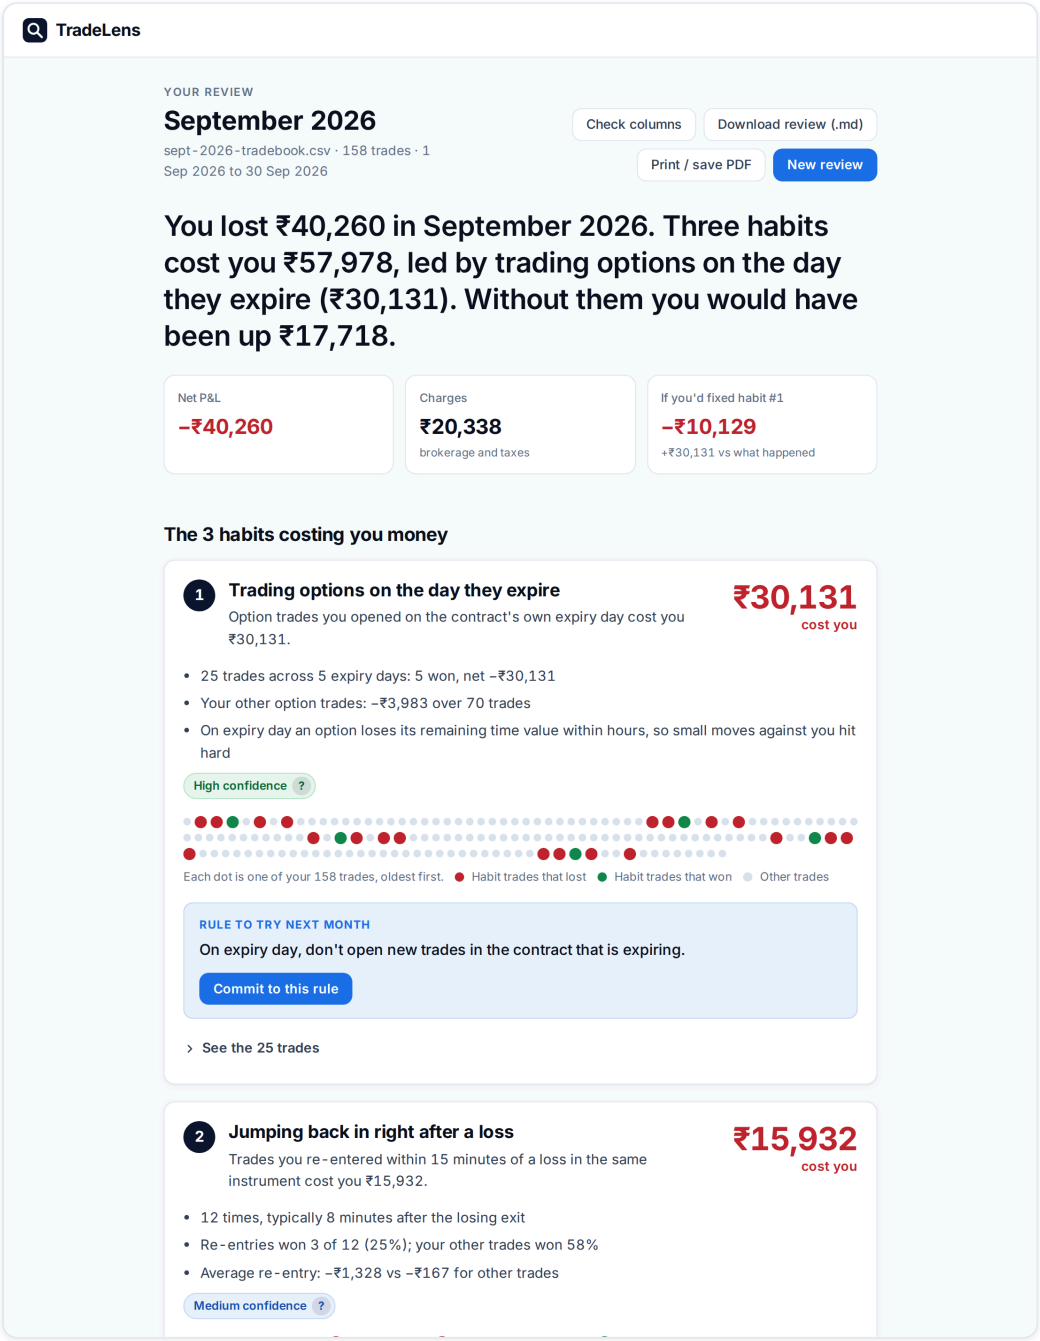
Limits: the data is synthetic; thresholds are judgement calls; only Zerodha-style tradebooks and generic closed-trade sheets are tested; charges are estimates when the file has none.

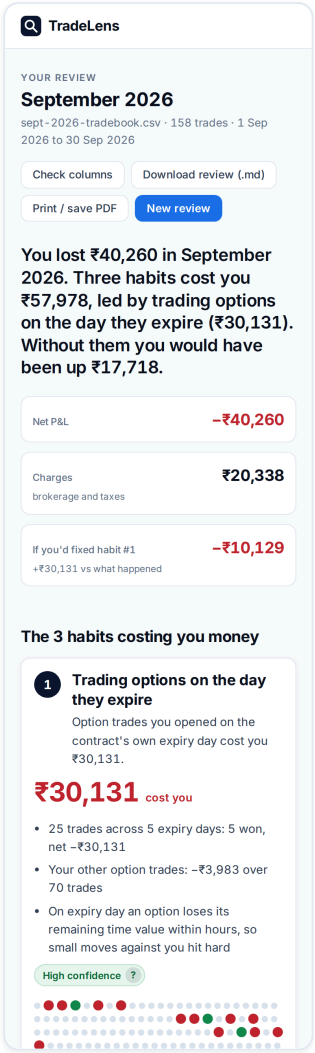
ARTIFACT 04 · WORKING APP

Click through it

No login, no backend. Files are analysed in the browser and never uploaded.

LIVE APP	trade-lens-omega.vercel.app
SOURCE	github.com/bhavyamohnot/TradeLens
FASTEST PATH	Open the app → Sample: a month of F&O order history → read the verdict and the three habit cards
ALSO TRY	"Sample: closed-trade sheet" (the brief's example file), or upload a Zerodha Console tradebook





What to look for

- **The verdict** names the month's net result, the habits' total cost and what the month would have been without them.
- **Each habit card** shows the ₹ cost, the evidence, a strip of the month's trades with the habit's trades highlighted, one rule, and "Commit to this rule".
- **"Checked, not costing you money"** shows the patterns that were tested and cleared, so the tool visibly doesn't invent problems.
- **Data notes** cover estimated charges and positions left open. The review downloads as Markdown or prints to PDF.